Manus AI: Product Purpose, Capabilities, Architecture, Target Users, Commercial Model, and Market Positioning

---

Purpose and Positioning

Manus is positioned as a next-generation, general-purpose autonomous AI agent that moves beyond suggestion-only chat interfaces to independently plan and execute complex, multi-step tasks across domains. This positioning is reflected in its mission to build "an autonomous AI agent that could handle any computer-based task," a mandate the company framed explicitly in its early roadmap communications (Manus Blog). "Three months ago, we launched Manus with a clear mission: to create an autonomous AI agent that could handle any computer-based task." [1] Public analyses consistently describe Manus as departing from conventional chatbots to independently plan, execute, and deliver results for multi-step tasks, underscoring its end-to-end orientation (BayTech Consulting). "Manus AI is designed not merely to respond or suggest, but to independently plan, execute, and deliver results for complex, multi-step tasks across various domains." [3]

The value proposition centers on business-grade autonomy, multi-modal understanding, deep tool integration, and transparency of execution. Independent reviewers highlight four differentiators: autonomous execution, multi-modal capabilities (text, images, code), advanced tool integration, and benchmark standing—features which Manus and third parties repeatedly emphasize (Single Grain). " *Autonomous Execution: Manus AI independently plans and completes multi-step tasks without continuous human prompting.*" [4] " **Multi-Modal Capabilities:** The system processes and generates text, images, and code across various applications." [4] " *Advanced Tool Integration: Manus seamlessly interacts with browsers, code editors, and other software to accomplish complex tasks.*" [4] " **Benchmark Performance:** Manus outperforms other AI agents on the GAIA benchmark for real-world problem-solving." [4] Its operating model is asynchronous and cloud-based, keeping work progressing when users are offline (BayTech

Consulting). “Tasks assigned to Manus AI can run in the background on cloud servers, meaning the process continues even if the user closes their browser, turns off their device, or is otherwise offline.” [3]

Manus’s origin story also shapes its positioning. It was built by Butterfly Effect Technology (maker of Monica), with an official launch on March 6, 2025 (Lawrence et al.). “Developed by China-based Butterfly Effect Technology (known for Monica) and officially launched on March 6, 2025, Manus AI represents a significant departure from conventional AI tools.” [24]

Core Features and Supported Use Cases

Manus emphasizes end-to-end autonomy in a managed cloud runtime, visible, iterative decision-making, and deep interaction with the web, code, files, and apps. Independent testing and vendor materials align on several pillars. Operationally, Manus runs independently in the cloud so tasks continue when users disconnect or go offline (Gocodeo; BayTech Consulting). “Manus AI operates independently in the cloud, which ensures that tasks continue to progress even if the initiating user disconnects.” [5] “Tasks assigned to Manus AI can run in the background on cloud servers, meaning the process continues even if the user closes their browser, turns off their device, or is otherwise offline.” [3] Within tasks, it actively browses and interacts with websites and surfaces its reasoning steps in real time (Gocodeo). “Manus AI actively browses the web, interacts with websites, and displays its decision-making steps in real time.” [5] It evolves performance through reinforcement learning signals during user interaction (Gocodeo). “Manus AI employs a reinforcement learning model that evolves through user interactions.” [5] It integrates beyond static datasets into dynamic sources such as X (Twitter), Telegram, and other APIs, expanding the scope of autonomous execution across live systems (Gocodeo). “Manus AI extends beyond static datasets by integrating with dynamic platforms such as X (formerly Twitter), Telegram, and other APIs.” [5]

Security and transparency of actions are addressed through per-session isolation and sandboxing and a visible “side-panel” of ongoing work. Independent testing describes a Linux sandbox, session isolation, and tool sandboxing to mitigate risks (DataCamp). “Linux sandbox environment: Manus operates within a controlled execution space, where it can install software, run scripts, and manipulate files.” [6] “Each Manus AI session operates in isolation, preventing users from accessing each other's execution environments.” [6] “Additionally, tools and commands are sandboxed, mitigating the risk of unauthorized system access.” [6] Analysts also point to the “Manus’s Computer” side panel that shows, in real time, what the agent is doing (BayTech Consulting). “This

component, often presented as a side panel in the user interface, provides real-time visibility into the AI's ongoing work." [3]

Use cases reflect breadth across research, analysis, content generation, planning, and operations, with many concrete illustrations from Manus's official gallery. These include personalized travel handbooks, teacher-ready video presentations, structured insurance comparisons, supplier scouting, YC database extraction, Amazon store analytics with visualizations, historical event visual maps, high-volume interview scheduling, and even building a custom teleprompter system (Manus Use Case Collection). "Manus integrates comprehensive travel information to create personalized itineraries and produces a custom travel handbook tailored specifically for your Japanese adventure." [11] "Manus develops engaging video presentations for middle school educators, clearly explaining the momentum theorem through accessible and educational content." [11] "Looking to compare insurance options? Manus generates clear, structured comparison tables highlighting key policy information with optimal recommendations tailored to your needs." [11] "Manus conducts comprehensive research across extensive networks to identify the most suitable suppliers for your specific requirements." [11] "Upload your Amazon store sales data and Manus delivers actionable insights, detailed visualizations, and customized strategies designed to increase your sales performance." [11] "Experience how Manus efficiently organizes interviews for 40 candidates with optimal time management to maximize productivity." [11] "Elevate your presentations with our customized teleprompter system featuring adjustable speed, scalable fonts, and enhanced readability for flawless delivery." [11]

Performance claims on the GAIA benchmark have featured prominently in coverage and community testing. External write-ups report state-of-the-art results across GAIA levels, with published scores and commentary suggesting superiority to competing agents and research systems (Gocodeo; Single Grain). "Manus AI has established new state-of-the-art (SOTA) performance across all three GAIA benchmark levels, outperforming OpenAI Deep Research and previous industry benchmarks." [5] " *Benchmark Performance:** Manus outperforms other AI agents on the GAIA benchmark for real-world problem-solving." [4]

Underlying AI Technologies and Technical Architecture

Manus's design is best understood as a hybrid of large language models (LLMs), a multi-agent coordination layer, and an executable action substrate, all running in a cloud

sandbox with integrated browser, shell, file system, and deployment capabilities. Technical investigations describe Manus as a wrapper over leading foundation models, especially Claude and Qwen families, augmented with retrieval and domain tooling (Renschni/GitHub Gist). “Manus is an autonomous AI agent built as a wrapper around foundation models (primarily Claude 3.5/3.7 and Alibaba's Qwen).” [7] “The developers confirmed Manus “supports retrieval augmented generation (RAG)”, meaning it combines external data retrieval with the model’s generation capabilities ([Manus在紅什麼？外媒評測訂餐、訂位、訂票…都碰壁：它是中國第二個DeepSeek時刻？]).” [7] Analysts corroborate the Claude/Qwen emphasis and suggest heavy integration rather than a proprietary base model (DataCamp). “Reports indicate the use of models such as Anthropic's Claude family (specifically mentioning Claude 3.5 Sonnet) and potentially fine-tuned versions of Alibaba's Qwen models.” [6] “* Over-reliance on existing models: Investigations into its architecture suggest that Manus heavily integrates Claude Sonnet and Qwen finetunes, rather than using a unique, proprietary model.” [6] Leanware also notes Manus’s dynamic model selection across Claude and Qwen based on task demands (Leanware). “Depending on the specific demands of each task, Manus integrates multiple AI models like Anthropic’s Claude 3.5 and Alibaba’s Qwen to perform these tasks.” [9]

A distinguishing design choice is the CodeAct-style execution model, where the agent emits and runs executable Python to take actions, tool calls, and even self-debug; integration with tools is standardized via function-call JSON or code until a final natural-language response is warranted (Renschni/GitHub Gist). “The system's key innovation is using executable Python code as its action mechanism (“CodeAct” approach), allowing it to perform complex operations autonomously.” [7] “Tool integration is standardized via a function-call interface: Manus outputs structured JSON or code specifying which tool to invoke; only upon task completion does it produce a natural-language response.” [7]

Multi-agent coordination and governance are embedded through a central executor and specialized sub-agents for planning, retrieval, generation, and verification, plus a strict, structured system prompt that encodes capability, browser, and coding rules (LinkedIn – Inside Manus AI; AI governance prompt, Gist). “A centralised "executor" agent collaborates with specialised subagents, each of whom is accountable for a distinct function, including planning, execution, and verification.” [8] “The prompt then lays out numerous **rules and guidelines** in a structured format (sections like `<system_capability>`, `<browser_rules>`, `<coding_rules>`, etc.) which the model must follow.” [7] Commentary also attributes enhanced chain-of-thought planning and action sequencing to a QwQ-32B component in some configurations (ThinkChina). “By integrating the QwQ-32B LLM, it achieves two key innovations: one, a “long chain-of-

thought reasoning" — instead of jumping straight to answers, the AI breaks down problems step-by-step (like showing your work in math class); two, action planning — it doesn't just think — it executes tasks in a logical order, much like following a recipe." [10]

At the infrastructure level, Manus runs in a cloud virtual compute environment with full access to a browser, shell, and code execution, plus a Python interpreter, Node.js libraries, terminal and file system access—even basic app deployment to public URLs (Renschni/Gist; LinkedIn; DataCamp). "It operates in a cloud-based virtual computing environment with full access to tools like web browsers, shell commands, and code execution." [7] "This environment contains a full autonomous web browser, a Python interpreter with libraries such as Node.js and Python, and terminal and file system access." [8] " *Integrated web browser control: Manus can navigate websites, extract data, interact with web elements, and even execute JavaScript within a browser console."* [6] " File system management: It can read, write, and organize files, making it useful for handling document-based workflows." [6] "* Deployment capabilities: Manus can deploy applications, including setting up websites and hosting services on public URLs." [6]

Target User Profiles and Industries

Manus is explicitly geared toward business users seeking to integrate AI autonomy without heavy custom-build costs, spanning both knowledge work and operational workflows (Spaculus). "Manus AI is designed specifically for businesses, making it a superior choice for companies looking to integrate AI into their operations without high customization costs." [2] Sector coverage is broad, with 2025-oriented examples across healthcare (diagnostics support, drug discovery, admin automation, predictive analytics), manufacturing (predictive maintenance, inventory, quality control, production automation), and retail (personalization, inventory optimization, customer insights, fraud detection) (Spaculus). "Industry-specific applications for 2025 include healthcare (AI-powered diagnostics, faster drug discovery, administrative automation, predictive analytics), manufacturing (predictive maintenance, inventory management, AI-based quality control, production automation), and retail (personalized shopping recommendations, inventory optimization, customer insights, fraud detection)." [2]

From a buyer-persona standpoint, adoption patterns often begin with entrepreneurs and marketers seeking time efficiency and creativity boosts, then expand into operations, finance, and IT where routine, complex business processes benefit from autonomous agents rather than one-off generative tasks (Jonathan Mast; INA & Associates). "Entrepreneurs and marketers prioritize time efficiency, making AI tools essential for automating tasks and boosting creativity." [18] "Generative AI is recommended for rapid,

creative, one-off tasks and industries like marketing, advertising, and education, while AI agents suit organizations with routine, complex business processes and data-driven strategies in manufacturing, finance, and real estate." [19] The official use case gallery further suggests relevance for research teams, sales operators, educators, HR coordinators, and SMB owners through the breadth of tasks Manus demonstrates (Manus Use Case Collection). "Manus expertly navigated the YC W25 database to identify all qualifying B2B companies, meticulously compiling this valuable information into a structured table." [11]

Pricing, Licensing, and Deployment Models

Manus uses a tiered, credit-based subscription with free and paid plans, and credit consumption tied to task complexity. External coverage and community breakdowns indicate five tiers: a Free plan with 1,000 starter credits and 300 daily refresh credits, a Basic plan at $19, a Plus plan at $39 with several thousand credits and Agent Mode, a Pro plan at $199 with significantly higher monthly credits and concurrency, and a Team plan at $39/seat with shared credit pools (Lindy). "Yes, Manus AI offers a free plan that gives you 1,000 starter credits and 300 credits refreshed daily." [13] "Manus AI costs $19 for the Basic plan." [13] "Plus | $39/month | Solo builders, creators | 3,900 credits/month (+3,900 promo, limited time), Agent Mode, 3 concurrent/scheduled tasks" [13] "Pro | $199/month | High-volume teams and power users | 19,900 credits/month (+promo), creative outputs, 10 concurrent/scheduled tasks" [13] "Team (still Beta) | $39/member/month, minimum 4 seats | Compliance, scale, or integrations | Shared pool of credits 3,900/member, exclusive data sources, early beta access" [13] Consumption and resets follow a usage-metered pattern, with credits consumed by task complexity and plan credits resetting monthly (Manus Docs). "Credits are consumed based on the complexity of the task." [15] "Plan credits reset monthly, but purchased add-on credits never expire and can be used anytime." [15] TechCrunch's launch coverage corroborates the paid tiers at $39 and $199 and notes concurrency and priority access in the higher tier (TechCrunch). "Manus AI, the viral AI agent platform out of China, on Monday morning launched two subscription plans starting at $39 per month." [12] "The other, costlier new plan, which costs $199 per month, grants users 19,900 credits, the ability to run five tasks simultaneously, and priority access during peak hours." [12]

Deployment today is SaaS-first with web, mobile, and desktop options, plus an API. The official site exposes feature pages, documentation, trust center, and API docs; it also advertises mobile and Windows apps (Manus site). "Resources BlogDocsUpdatesHelp centerTrust centerAPITeam planStartupsPlaybookBrand assets" [16] "Download Mobile

appWindows appMy Browser” [16] Current access has been staged through a closed or invite-only beta with paid plans added shortly after launch (DataCamp; TechCrunch). “Access to Manus AI is currently limited to an invitation-only beta phase.” [6] “Manus AI, the viral AI agent platform out of China, on Monday morning launched two subscription plans starting at $39 per month.” [12]

Licensing terms specific to Manus.im’s SaaS were not found in the public materials reviewed; however, the presence of API documentation and team plans suggests a standard SaaS subscription framework with usage metering and account provisioning (Manus site). “Resources BlogDocsUpdatesHelp centerTrust centerAPITeam planStartupsPlaybookBrand assets” [16]

Comparative Positioning vs. Similar AI Solutions

Autonomous agent platforms span from managed, end-user services like Manus to open-source frameworks and builder platforms (e.g., AutoGPT, LangChain) that require more engineering but offer greater hosting and customization control.

AutoGPT exemplifies the open-source path. It allows users to create, deploy, and manage autonomous agents locally or in their own cloud accounts and is free to use, with a low-code agent builder, workflow management, and multi-LLM integrations (IBM; DataCamp Tutorial; Futurepedia). “AutoGPT is an open-source [artificial intelligence (AI)] platform that allows users to automate multistep projects and complex workflows with [AI agents] based on OpenAI’s GPT-4 [large language model (LLM)].” [18] “Available on GitHub, this platform offers a comprehensive suite of tools to create, deploy, and manage AI agents that automate complex workflows.” [20] “It is pre-integrated with LLM providers such as OpenAI, Anthropic, Groq, and Llama so that a wide range of tasks like data processing, content creation, or even fun activities can be automated.” [19]

LangChain, by contrast, is a developer platform and production toolkit for building agents, offering Cloud, Hybrid, or Self-Hosted deployment options, seat-based pricing for observability and deployment, and primitives (LangGraph) for custom workflows (LangChain Pricing; LangChain Official). “Cloud, Hybrid, or Self-Hosted” [21] “$0 / seat per month” [21] “$39 / seat per month” [21] “LangChain helps you ship quickly with less code using a pre-built agent architecture and model integrations.” [22] “LangGraph puts you in control with low-level primitives to build custom agent workflows.” [22]

Manus differentiates as a managed, end-user agent with cloud sandboxes, visible autonomy, browser/file/shell access, and credit-metered pricing—all packaged for

business users without requiring in-house agent engineering (Gocodeo; Renschni/Gist). “Manus AI operates independently in the cloud, which ensures that tasks continue to progress even if the initiating user disconnects.” [5] “It operates in a cloud-based virtual computing environment with full access to tools like web browsers, shell commands, and code execution.” [7] “Tool integration is standardized via a function-call interface: Manus outputs structured JSON or code specifying which tool to invoke; only upon task completion does it produce a natural-language response.” [7]

Manus highlights SOTA GAIA performance and advanced parallel agent orchestration (Wide Research), enabling large-scale analysis. “Manus AI has established new state-of-the-art (SOTA) performance across all three GAIA benchmark levels, outperforming OpenAI Deep Research and previous industry benchmarks.” [5] “Wide Research: Distributes tasks across multiple agents for parallel processing” [17] “Manus AI has introduced an exciting new feature called Wide Research, which allows users to use more than 100 AI agents at the same time.” [17]

Product Maturity, Roadmap, and Known Limitations

As an emerging agent system, Manus’s rollout and product roadmap indicate rapid iteration and expansion of integrations. The vendor reports introducing a Chat Mode for daily tasks, Playbooks for onboarding, doubling speeds and cutting costs fivefold through architectural/infrastructure changes, with planned features like Scheduled Tasks, Wide Research expansion, Omni Search, and broader app integrations including email, calendar, drive, and task tools (Manus Blog). “First, we introduced Chat Mode for more efficient daily tasks.” [1] “Second, we developed comprehensive Playbooks for new users to get started.” [1] “Over the past three months, we’ve increased speeds 2x and brought costs down 5x through architectural enhancements and infrastructure optimization.” [1] “We're developing Scheduled Tasks for recurring research, Wide Research capabilities that span diverse sources, Omni Search that understands intent and context, and access to more specialized domain data sources.” [1] “We're integrating into more daily working apps to streamline workflows, including Email management and composition, Calendar preparation and follow-up, Drive organization and File management, and Task management coordination across projects and teams.” [1]

Independent reviews also flag current limitations typical of early-stage agents. Access is invite-only; some users report loops or getting stuck when tasks are ill-defined; and architecture investigations point to heavy reliance on Claude and Qwen rather than a proprietary base model (DataCamp). “Access to Manus AI is currently limited to an

invitation-only beta phase." [6] " *Glitches and inconsistencies: Some users have reported that Manus loops or gets stuck in repetitive cycles, struggling with complex decision-making when tasks aren't well-defined."* [6] " Over-reliance on existing models: Investigations into its architecture suggest that Manus heavily integrates Claude Sonnet and Qwen finetunes, rather than using a unique, proprietary model." [6]

Capability-to-Use Case-to-Impact Mapping

Manus's most distinctive capabilities—autonomous cloud execution, active web interaction with transparent reasoning, sandboxed code/shell/browser control, multi-model orchestration, and parallel "Wide Research"—map to enterprise use cases with measurable operational implications.

- Autonomy and asynchronous execution reduce supervision and enable long-running tasks such as market landscaping, supplier discovery, or data processing to complete without constant human presence. These behaviors are directly supported by the agent's cloud background execution design. "Tasks assigned to Manus AI can run in the background on cloud servers, meaning the process continues even if the user closes their browser, turns off their device, or is otherwise offline." [3]

- Active web interaction and a visible side panel make Manus suitable for workflows requiring browser automation, source verification, and auditability, such as regulatory tracking, competitive scans, and structured comparisons (e.g., insurance), which Manus demonstrates. "Manus AI actively browses the web, interacts with websites, and displays its decision-making steps in real time." [5] "This component, often presented as a side panel in the user interface, provides real-time visibility into the AI's ongoing work." [3] "Looking to compare insurance options? Manus generates clear, structured comparison tables highlighting key policy information with optimal recommendations tailored to your needs." [11]

- Code/shell/file system access and deployment abilities extend Manus beyond "chat" into application scaffolding, data wrangling, and light DevOps, supporting use cases like report generation with visualizations, data cleaning, and publishing outputs to accessible endpoints. " *File system management: It can read, write, and organize files, making it useful for handling document-based workflows."* [6] " Deployment capabilities: Manus can deploy applications, including setting up websites and hosting services on public URLs." [6]

- Multi-model orchestration and planning yield better robustness across varied tasks—e.g., switching between reasoning, browsing, coding, and retrieval—while Wide Research supports large-scale parallel literature or market reviews. “Depending on the specific demands of each task, Manus integrates multiple AI models like Anthropic’s Claude 3.5 and Alibaba’s Qwen to perform these tasks.” [9] “Manus AI has introduced an exciting new feature called Wide Research, which allows users to use more than 100 AI agents at the same time.” [17]

These linkages suggest Manus’s highest ROI in scenarios where workflows are multi-step, data-rich, and require both web/tool interaction and durable execution windows—exactly where traditional chat-based assistants struggle to maintain context, continuity, and actionability.

Conclusion and Buyer’s Takeaways

Manus combines LLM-driven reasoning, a multi-agent planning and governance layer, and an executable code/browser/shell environment to deliver end-to-end autonomous task completion for business users. Its strengths are clear: asynchronous cloud execution with visible autonomy, deep tool use, and a growing portfolio of integrations and parallel research capabilities; early benchmarks and use case demonstrations underscore real-world viability. At the same time, pragmatic buyers should account for early-stage limitations (invite-only access, occasional looping on ill-posed tasks, and reliance on third-party base models). For teams deciding among Manus, open-source frameworks like AutoGPT, or builder platforms like LangChain, the primary trade-off is between time-to-value and control: Manus maximizes immediate utility for non-developer teams, while AutoGPT/LangChain maximize customization and hosting flexibility.

Decision lens	Manus (managed)	AutoGPT (open source)	LangChain (builder)
Time to value	Fastest for non-developers; turnkey SaaS	Moderate; DIY setup and ops	Moderate to slow; engineering-led
Control & hosting	Low–medium; vendor cloud	High; user-controlled infra	High; supports self-hosted
Observability		Varies by setup	

Decision lens	Manus (managed)	AutoGPT (open source)	LangChain (builder)
	Visible “computer” side panel		Full observability stack
Scaling research	“Wide Research” (100+ agents)	Parallelism requires DIY	Orchestrate via LangGraph
When to choose	Business users needing autonomous execution now	Teams needing control, zero license fees	Product/ML teams building custom agents

Across industries like healthcare, manufacturing, retail, and finance, this capability stack translates into accelerated research, structured decision support, and automation of routine digital operations. Selecting Manus versus alternatives hinges on whether a managed, credit-based agent with strong web/tool action is more valuable than the flexibility and ownership of building and hosting custom agents.

Sources

[1] Manus.im Blog – What we saw in the past three months. And what we see in the future: https://manus.im/blog/what-we-saw-in-the-past-three-months-and-what-we-see-in-the-future
[2] Spaculus – Manus AI: The Future of Business AI or Just Another AI Agent?: https://spaculus.com/news/manus-ai-business-automation-guide/
[3] BayTech Consulting – Manus AI: An Analytical Guide to the Autonomous AI Agent 2025: https://www.baytechconsulting.com/blog/manus-ai-an-analytical-guide-to-the-autonomous-ai-agent-2025
[4] Single Grain – Manus AI: The Ultimate Guide to Understanding and Using It: https://www.singlegrain.com/digital-marketing/manus-ai-the-ultimate-guide-to-understanding-and-using-it/
[5] Gocodeo – Manus AI: Capabilities, GAIA Benchmark, Use Cases & More: https://www.gocodeo.com/post/manus-ai-capabilities
[6] DataCamp – Manus AI: Features, Architecture, Access, Early Issues & More: https://www.datacamp.com/blog/manus-ai
[7] Renschni (GitHub Gist) – In-depth technical investigation into the Manus AI agent: https://gist.github.com/renschni/4fbc70b31bad8dd57f3370239dccd58f
[8] LinkedIn – Edition 9: Inside Manus AI – Architecture and Benchmarks: https://www.linkedin.com/pulse/edition-9-inside-manus-ai-architecture-benchmarks-futureagi-

vd7ee
[9] Leanware – Manus AI Agent: What It Is, How It Works, & Its Impact [2025]: https://www.leanware.co/insights/manus-ai-agent
[10] ThinkChina – Manus: Is China's overnight AI sensation a breakthrough or just hype?: https://www.thinkchina.sg/technology/manus-chinas-overnight-ai-sensation-breakthrough-or-just-hype
[11] Manus – Official use case collection: https://manus.im/usecase-official-collection
[12] TechCrunch – Manus launches paid subscription plans and a mobile app: https://techcrunch.com/2025/03/31/manus-launches-paid-subscription-plans-and-a-mobile-app/
[13] Lindy – Manus AI Pricing: A Detailed Breakdown of Each Plan: https://www.lindy.ai/blog/manus-ai-pricing
[14] Manus – Pricing: https://manus.im/pricing
[15] Manus Docs – Plans and Pricing: https://manus.im/docs/introduction/plans
[16] Manus Website – Product, resources, downloads: https://manus.im/
[17] LinkedIn – Manus AI Wide Research: Deploy 100+ Agents for AI Automation: https://www.linkedin.com/pulse/manus-ai-launches-wide-research-deploying-100-agents-nantha-kumar-l-ag3fc
[18] IBM – What is AutoGPT?: https://www.ibm.com/think/topics/autogpt
[19] DataCamp Tutorial – AutoGPT Guide: Creating And Deploying Autonomous AI Agents Locally: https://www.datacamp.com/tutorial/autogpt-guide
[20] Futurepedia – AutoGPT AI Reviews: Use Cases, Pricing & Alternatives: https://www.futurepedia.io/tool/auto-gpt
[21] LangChain – Plans and Pricing: https://www.langchain.com/pricing
[22] LangChain – Official Site: https://www.langchain.com/
[24] Lawrence et al. – Beyond Chatbots: Understanding Manus AI: https://lawrence.eti.br/2025/04/29/beyond-chatbots-understanding-manus-ai-the-general-ai-agent-changing-everything/